

Cryptography

Michele Laurenti

January 5, 2017

Abstract

Notes taken during the Cryptography lectures held by Daniele Venturi (<http://danieleventuri.altervista.org/crypto.shtml>) in fall 2016 at Sapienza.

Contents

1	Introduction	1
1.1	Perfect secrecy	1
1.2	Authentication	4
1.3	Randomness Extraction	6
2	Computational Cryptography	7
2.1	One Way Functions	7
2.2	Computational Indistinguishability	8
2.3	Pseudo Random Generators	9
	Acronyms	11

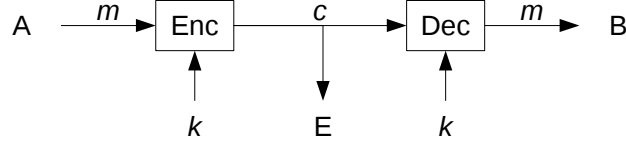


Figure 1: Message exchange between A and B using symmetric encryption. E is the eavesdropper.

1 Introduction

Solomon,
I'm concerned about security; I think, when we email each other,
we should use some sort of code.

Confidentiality is our goal. We want to encrypt and decrypt a (plaintext) message m , using a key, to obtain a cyphertext c . As per Kirkoff's principle, only the key is secret.

Our encryption schemes have the following syntax:

$$\Pi = (\text{Gen}, \text{Enc}, \text{Dec}).$$

A and B , the actors of our communication exchange (fig. 1), share k , the key, taken from some key space \mathcal{K} . The elements of our encryption scheme play the following roles:

1. Gen outputs a random key from the key space \mathcal{K} , and we write this as $k \leftarrow \$\text{Gen}$;
2. $\text{Enc} : \mathcal{K} \times \mathcal{M} \rightarrow \mathcal{C}$ is the encryption function, mapping a key and a message to a cyphertext;
3. $\text{Dec} : \mathcal{K} \times \mathcal{C} \rightarrow \mathcal{M}$ is the decryption function, mapping a key and a cyphertext to a message.

We expect an encryption scheme to be at least correct:

$$\forall k \in \mathcal{K}, \forall m \in \mathcal{M}. \text{Dec}(k, \text{Enc}(k, m)) = m.$$

1.1 Perfect secrecy

Shannon defined “perfect secrecy”, *i.e.*, the fact that the cyphertext carries no information about the plaintext.

Definition 1. [Perfect secrecy] Let M be a random variable (RV) over \mathcal{M} , and K be a uniform distribution over \mathcal{K} .

(Enc, Dec) has perfect secrecy if

$$\forall M, \forall m \in \mathcal{M}, c \in \mathcal{C}. \Pr [M = m] = \Pr [M = m | C = c]$$

where $C = \text{Enc}(k, m)$ is a third RV. \diamond

We have equivalent definitions for perfect secrecy.

Theorem 1. *The following definitions are equivalent:*

1. *definition 1;*
2. *M and C are independent;*
3. *$\forall m, m' \in \mathcal{M}, \forall c \in \mathcal{C}$*

$$\Pr [\text{Enc}(k, m) = c] = \Pr [\text{Enc}(k, m') = c]$$

where k is a random key in \mathcal{K} chosen with uniform probability. \diamond

Proof of theorem 1. First, we show that 1 implies 2.

$$\begin{aligned} \Pr [M = m] &= \Pr [M = m | C = c] \\ &= \frac{\Pr [M = m \wedge C = c]}{\Pr [C = c]} && \text{(by Bayes)} \\ &\implies \\ \Pr [M = m] \Pr [C = c] &= \Pr [M = m \wedge C = c] \end{aligned}$$

which is the definition of independence.

Now we show that 2 implies 3. Fix $m \in \mathcal{M}$ and $c \in \mathcal{C}$.

$$\begin{aligned} \Pr [\text{Enc}(k, m) = c] &= \Pr [\text{Enc}(k, M) = c | M = m] && \text{(we fixed } m) \\ &= \Pr [C = c | M = m] && \text{(definition of the RV } C) \\ &= \Pr [C = c]. && \text{(by 2)} \end{aligned}$$

Since m is arbitrary, we can do the same for m' , and obtain

$$\Pr [\text{Enc}(k, m') = c] = \Pr [C = c]$$

which gives us 3.

Now we want to show that 3 implies 1. Take any $c \in \mathcal{C}$.

$$\begin{aligned}
\Pr[C = c] &= \sum_{m' \in \mathcal{M}} \Pr[C = c \wedge M = m'] \\
&= \sum_{m' \in \mathcal{M}} \Pr[C = c | M = m'] \Pr[M = m'] && \text{(by Bayes)} \\
&= \sum_{m' \in \mathcal{M}} \Pr[\text{Enc}(k, M) = c | M = m'] \Pr[M = m'] \\
&= \sum_{m' \in \mathcal{M}} \Pr[\text{Enc}(k, m') = c] \Pr[M = m'] \\
&= \Pr[\text{Enc}(k, m) = c] \underbrace{\sum_{m' \in \mathcal{M}} \Pr[M = m']}_1 \\
&\quad \text{(Enc is independent of } M, \text{ so we take it out)} \\
&= \Pr[\text{Enc}(k, M) = c | M = m] = \Pr[C = c | M = m].
\end{aligned}$$

We are left to show that $\Pr[M = m] = \Pr[M = m | C = c]$, but this is easy with Bayes. \square

One Time Pad

Now we'll see a perfect encryption scheme, the One Time Pad (OTP). The message space, the cyphertext space, and the key space are all the same, *i.e.*, $\mathcal{M} = \mathcal{K} = \mathcal{C} = \{0, 1\}^l$, with $l \in \mathbb{N}^+$.

Encryption and decryption use the xor operation:

- $\text{Enc}(k, m) = k \oplus m = c$;
- $\text{Dec}(k, c) = c \oplus k = (k \oplus m) \oplus k = m$.

Seeing that this is correct is immediate.

This can actually be done in any finite abelian group $(\mathbb{G}, +)$, where you just do $k + m$ to encode and $c - k$ to decode.

Theorem 2. *OTP is perfectly secure.* \diamond

Proof of theorem 2. Fix $m \in \mathcal{M}, c \in \mathcal{C}$, and choose a random key.

$$\Pr[\text{Enc}(k, m) = c] = \Pr[k = c - m] = \frac{1}{|\mathbb{G}|}.$$

This is true for any m , so we are done. \square

OTP has two problems:

1. the key is long (as long as the message);

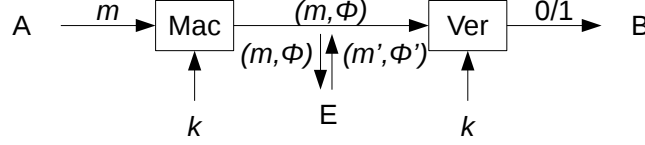


Figure 2: Message exchange between A and B using symmetric authentication. E is the eavesdropper.

2. we can't reuse the key:

$$\begin{aligned} c &= k + m \\ c' &= k + m' \end{aligned} \implies c - c' = m - m' \implies m' = m - (c - c').$$

Theorem 3. [Shannon, 1949] In any perfectly secure encryption scheme the size of the key space is at least as large as the size of the message space, i.e., $|\mathcal{K}| \geq |\mathcal{M}|$. \diamond

Proof of theorem 3. Assume, for the sake of contradiction, that $|\mathcal{K}| < |\mathcal{M}|$. Fix M to be the uniform distribution over \mathcal{M} , which we can do as perfect secrecy works for any distribution. Take a cyphertext $c \in \mathcal{C}$ such that $\Pr[C = c] > 0$, i.e., $\exists m, k$ such that $\text{Enc}(k, m) = c$.

Consider $\mathcal{M}' = \{\text{Dec}(k, c) : k \in \mathcal{K}\}$, the set of all messages decrypted from c using any key. Clearly, $|\mathcal{M}'| \leq |\mathcal{K}| < |\mathcal{M}|$, so $\exists m' \in \mathcal{M}$ such that $m' \notin \mathcal{M}'$. This means that

$$\Pr[M = m'] = \frac{1}{|\mathcal{M}|} \neq \Pr[M = m' | C = c] = 0$$

in contradiction with perfect secrecy. \square

In the rest of the course we will forget about perfect secrecy, and strive for computational security, i.e., bound the computational power of the adversary.

1.2 Authentication

The aim of authentication is to avoid tampering of E with the messages exchanged between A and B (fig. 2).

A Message Authentication Code (MAC) is defined as a tuple $\Pi = (\text{Gen}, \text{Mac}, \text{Vrfy})$, where:

- Gen , as usual, outputs a random key from some key space \mathcal{K} ;
- $\text{Mac} : \mathcal{K} \times \mathcal{M} \rightarrow \Phi$ maps a key and a message to an authenticator in some authenticator space Φ ;
- $\text{Vrfy} : \mathcal{K} \times \mathcal{M} \times \Phi \rightarrow \{0, 1\}$ verifies the authenticator.

As usual, we expect a MAC to be correct, *i.e.*,

$$\forall m \in \mathcal{M}, \forall k \in \mathcal{K}. \text{Vrfy}(k, m, \text{Mac}(k, m)) = 1.$$

If the Mac function is deterministic, then it must be that $\text{Vrfy}(k, m, \phi) = 1$ if and only if $\text{Mac}(k, m) = \phi$.

Security for MACs is that *forgery* must be hard: you can't come up with an authenticator for a message if you don't know the key.

Definition 2. [Information theoretic MAC] $(\text{Mac}, \text{Vrfy})$ has ε -statistical security if for all (possibly unbounded) adversary \mathcal{A} , for all $m \in \mathcal{M}$,

$$\Pr \left[\begin{array}{l} \text{Vrfy}(k, m', \phi') = 1 \wedge m' \neq m : \\ \begin{array}{l} k \leftarrow \text{sKeyGen}; \\ \phi = \text{Mac}(k, m); \\ (m', \phi') \leftarrow \mathcal{A}(m, \phi) \end{array} \end{array} \right] \leq \varepsilon$$

i.e., the adversary forges a (m', ϕ') that verifies with key k with low probability, even if it knows a valid pair (m, ϕ) . \diamond

As an exercise, prove that the above is impossible if $\varepsilon = 0$.

Information theoretic security is also called unconditional security. Later we'll see *conditional* security, based on computational assumptions.

Definition 3. [Pairwise independence] Given a family $\mathcal{H} = \{h_k : \mathcal{M} \rightarrow \Phi\}_{k \in \mathcal{K}}$ of functions, we say that \mathcal{H} is pairwise independent if for all distinct m, m' we have that $(h_k(m), h_k(m')) \in \Phi^2$ is uniform over the choice of $k \leftarrow \mathcal{K}$. \diamond

We say straight away a construction of a pairwise independent family of function. Let p be a prime, the functions in our family will be

$$h_{a,b}(m) = am + b \pmod{p}$$

with $\mathcal{K} = \mathbb{Z}_p^2$, and with $\mathcal{M} = \Phi = \mathbb{Z}_p$.

Theorem 4. *The above construction is pairwise independent.* \diamond

Proof of theorem 4. For any m, m', ϕ, ϕ' , we want to find the value of

$$\Pr [am + b = \phi \wedge am' + b = \phi']$$

for $a, b \leftarrow \mathbb{Z}_p^2$. This is the same as

$$\Pr_{a,b} \left[\begin{pmatrix} m & 1 \\ m' & 1 \end{pmatrix} \begin{pmatrix} a \\ b \end{pmatrix} = \begin{pmatrix} \phi \\ \phi' \end{pmatrix} \right] = \Pr_{a,b} \left[\begin{pmatrix} a \\ b \end{pmatrix} = \begin{pmatrix} m & 1 \\ m' & 1 \end{pmatrix}^{-1} \begin{pmatrix} \phi \\ \phi' \end{pmatrix} \right] = \frac{1}{|\Phi|^2}.$$

This is true since $\begin{pmatrix} m & 1 \\ m' & 1 \end{pmatrix}^{-1} \begin{pmatrix} \phi \\ \phi' \end{pmatrix}$ is just a couple of (constant) numbers, so the probability of choosing (a, b) such that they equal the constant is just $\frac{1}{|\Phi|^2}$. \square

If h_k is part of a pairwise independent family of functions, then $\text{Mac}(k, m) = h_k(m)$, and $\text{Vrfy}(k, m, \phi)$ is simply computing $h_k(m)$ and comparing it with ϕ , *i.e.*,

$$\text{Vrfy}(k, m, \phi) = 1 \iff h_k(m) = \phi.$$

We now prove that this is an information theoretic MAC.

Theorem 5. *Any pairwise independent function is $\frac{1}{|\Phi|}$ -statistical secure.* \diamond

Proof of theorem 5. Take any two distinct m, m' , and two ϕ, ϕ' . We show that the probability that $\text{Mac}(k, m') = \phi'$ is exponentially small.

$$\Pr_k [\text{Mac}(k, m) = \phi] = \Pr_k [h_k(m) = \phi] = \frac{1}{|\Phi|}.$$

Now look at the joint probabilities:

$$\begin{aligned} \Pr_k [\text{Mac}(k, m) = \phi \wedge \text{Mac}(k, m') = \phi'] &= \Pr_k [h_k(m) = \phi \wedge h_k(m') = \phi'] \\ &\quad \text{(by definition)} \\ &= \frac{1}{|\Phi|^2} = \frac{1}{|\Phi|} \cdot \frac{1}{|\Phi|}. \end{aligned}$$

The last steps come from the fact that h_k is pairwise independent. To see that the construction is $\frac{1}{|\Phi|}$ -statistical secure:

$$\begin{aligned} \Pr_k [\text{Mac}(k, m') = \phi' | \text{Mac}(k, m) = \phi] &= \Pr_k [h_k(m') = \phi' | h_k(m) = \phi] \\ &= \frac{\Pr_k [h_k(m) = \phi \wedge h_k(m') = \phi']}{\Pr_k [h_k(m) = \phi]} \\ &= \frac{1}{|\Phi|}. \end{aligned}$$

□

Note that the previous construction ($h_k(m) = am + b \pmod p$) is insecure if the same key $k = (a, b)$ is used for two messages.

Theorem 6. *Any t -time $2^{-\lambda}$ -statistically secure MAC has key of size $(t + 1)\lambda$.* \diamond

1.3 Randomness Extraction

X is a random source (possibly not uniform). $\text{Ext}(X) = Y$ is a uniform RV.

First, let's see a construction for a binary RV. Let B be a RV such that $\Pr[B = 1] = p$ and $\Pr[B = 0] = 1 - p$, with $p \neq 1 - p$. We take two samples, B_1 and B_2 from B , and we want to obtain an unbiased RV B' .

1. Take two samples, $b_1 \leftarrow \$B_1$ and $b_2 \leftarrow \$B_2$;
2. if $b_1 = b_2$, sample again;
3. if $(b_1 = 1 \wedge b_2 = 0)$, output 1; if $(b_1 = 0 \wedge b_2 = 1)$, output 0.

It's easy to verify that B' is uniform:

$$\begin{aligned}\Pr[B' = 1] &= \Pr[B_1 = 1 \wedge B_2 = 0] = p(1 - p) \\ \Pr[B' = 0] &= \Pr[B_1 = 0 \wedge B_2 = 1] = (1 - p)p.\end{aligned}$$

How many trials do we have to make before outputting something? $2(1 - p)p$ is the probability that we output something. The probability that we don't output anything for n steps is thus $(1 - 2(1 - p)p)^n$.

2 Computational Cryptography

To introduce computational cryptography we first have to define a computational model. We assume the adversary is efficient, *i.e.*, it is a Probabilistic Polynomial Time (PPT) adversary.

We want that the probability of success of the adversary is tiny, *i.e.*, negligible for some $\lambda \in \mathbb{N}$. A function $\varepsilon : \mathbb{N} \rightarrow \mathbb{R}$ is negligible if $\forall c > 0. \exists n_0$ such that $\forall n > n_0. \varepsilon(n) < n^{-c}$.

We rely on computational assumptions, *i.e.*, in tasks believed to be hard for any efficient adversary. In this setting we make conditional statements, *i.e.*, if a certain assumption holds then a certain crypto-scheme is secure.

2.1 One Way Functions

A simple computational assumption is the existence of One Way Functions (OWFs), *i.e.*, functions for which is hard to compute the inverse.

Definition 4. [One Way Function] A function $f : \{0, 1\}^* \rightarrow \{0, 1\}^*$ is a OWF if $f(x)$ can be computed in polynomial time for all x and for all PPT adversaries \mathcal{A} it holds that

$$\Pr[f(x') = y : x \leftarrow \$\{0, 1\}^*; y = f(x); x' \leftarrow \mathcal{A}(1^\lambda, y)] \leq \varepsilon(\lambda). \quad \diamond$$

The 1^λ given to the adversary \mathcal{A} is there to highlight the fact that \mathcal{A} is polynomial in the length of the input (λ).

Russel Impagliazzo proved that OWFs are equivalent to One Way Puzzles, *i.e.*, couples (Pgen, Pver) where $\text{Pgen}(1^\lambda) \rightarrow (y, x)$ gives us a puzzle (y) and a solution to it (x), while $\text{Pver}(x, y) \rightarrow 0/1$ verifies if x is a solution of y .

Another object of interest in this classification are average hard NP-puzzles, for which you can only get an instance, *i.e.*, $\text{Pgen}(1^\lambda) \rightarrow y$.

Impagliazzo says we live in one of five worlds:

1. Algorithmica, where $P = NP$;
2. Heuristica, where there are no average hard NP-puzzles, *i.e.*, problems without solution;
3. Pessiland, where you have average hard NP-puzzles;
4. Minicrypt, where you have OWF, one-way NP-puzzles, but no Public Key Cryptography (PKC);
5. Cryptomania, where you have both OWF and PKC.

We'll stay in Minicrypt for now.

OWF are hard to invert on average. Two examples:

- factoring the product of two large prime numbers;
- compute the discrete logarithm, *i.e.*, take a finite group (\mathbb{G}, \cdot) , and compute $y = g^x$ for some $g \in \mathbb{G}$. The find $x = \log_g(y)$. This is hard to compute in some groups, *e.g.*, \mathbb{Z}_p^* .

2.2 Computational Indistinguishability

Definition 5. [Distribution Ensemble] A distribution ensemble $\mathcal{X} = \{X_n\}_{n \in \mathbb{N}}$ is a sequence of distributions X_i over some space $\{0, 1\}^\lambda$. \diamond

Definition 6. [Computational Indistinguishability] Two distribution ensembles \mathcal{X}_λ and \mathcal{Y}_λ are computationally indistinguishable, written as $\mathcal{X}_\lambda \approx_c \mathcal{Y}_\lambda$, if for all PPT distinguishers \mathcal{D} it holds that

$$\left| \Pr[\mathcal{D}(\mathcal{X}_\lambda) = 1] - \Pr[\mathcal{D}(\mathcal{Y}_\lambda) = 1] \right| \leq \varepsilon(\lambda).$$

\diamond

Lemma 1. [Reduction] If $\mathcal{X} \approx_c \mathcal{Y}$, then for all PPT functions f , $f(\mathcal{X}) \approx_c f(\mathcal{Y})$. \diamond

Proof of lemma 1. Assume, for the sake of contradiction, that $\exists f$ such that $f(\mathcal{X}) \not\approx_c f(\mathcal{Y})$: then we can distinguish \mathcal{X} and \mathcal{Y} . Since $f(\mathcal{X}) \not\approx_c f(\mathcal{Y})$, then $\exists p = \text{poly}(\lambda), \mathcal{D}$ such that, for infinitely many λ s

$$\left| \Pr[\mathcal{D}(f(\mathcal{X}_\lambda)) = 1] - \Pr[\mathcal{D}(f(\mathcal{Y}_\lambda)) = 1] \right| \geq \frac{1}{p(\lambda)}.$$

\mathcal{D} distinguishes \mathcal{X}_λ and \mathcal{Y}_λ with non-negligible probability. Consider the following \mathcal{D}' , which is given

$$z = \begin{cases} x \leftarrow \mathcal{X}_\lambda; \\ y \leftarrow \mathcal{Y}_\lambda. \end{cases}$$

\mathcal{D}' runs $\mathcal{D}(f(z))$ and outputs whatever it outputs, and has the same probability of distinguishing \mathcal{X} and \mathcal{Y} of \mathcal{D} , in contradiction with the fact that $\mathcal{X} \approx_c \mathcal{Y}$. \square

Now we show that computational indistinguishability is transitive.

Lemma 2. *[Hybrid Argument] Let $\mathcal{X} = \{X_\lambda\}$, $\mathcal{Y} = \{Y_\lambda\}$, $\mathcal{Z} = \{Z_\lambda\}$ be distribution ensembles. If $\mathcal{X}_\lambda \approx_c \mathcal{Y}_\lambda$ and $\mathcal{Y}_\lambda \approx_c \mathcal{Z}_\lambda$, then $\mathcal{X}_\lambda \approx_c \mathcal{Z}_\lambda$. \diamond*

Proof of lemma 2. This follows from the triangular inequality.

$$\begin{aligned}
\left| \Pr[\mathcal{D}(\mathcal{X}_\lambda) = 1] - \Pr[\mathcal{D}(\mathcal{Z}_\lambda) = 1] \right| &= \left| \Pr[\mathcal{D}(\mathcal{X}_\lambda) = 1] - \Pr[\mathcal{D}(\mathcal{Y}_\lambda) = 1] \right. \\
&\quad \left. + \Pr[\mathcal{D}(\mathcal{Y}_\lambda) = 1] - \Pr[\mathcal{D}(\mathcal{Z}_\lambda) = 1] \right| \\
&\leq \left| \Pr[\mathcal{D}(\mathcal{X}_\lambda) = 1] - \Pr[\mathcal{D}(\mathcal{Y}_\lambda) = 1] \right| \\
&\quad + \left| \Pr[\mathcal{D}(\mathcal{Y}_\lambda) = 1] - \Pr[\mathcal{D}(\mathcal{Z}_\lambda) = 1] \right| \\
&\leq 2\varepsilon(\lambda). \quad (\text{negligible})
\end{aligned}$$

□

We often prove $\mathcal{X} \approx_c \mathcal{Y}$ by defining a sequence $\mathcal{H}_0, \mathcal{H}_1, \dots, \mathcal{H}_t$ of distributions ensembles such that $\mathcal{H}_0 \equiv \mathcal{X}$ and $\mathcal{H}_t \equiv \mathcal{Y}$, and that for all i , $\mathcal{H}_i \approx_c \mathcal{H}_{i+1}$.

2.3 Pseudo Random Generators

Let's see our first cryptographic primitive. Pseudo Random Generators (PRGs) take in input a random seed and generate pseudo random sequences with some stretch, *i.e.*, output longer than input, and indistinguishable from a true random sequence.

Definition 7. [Pseudo Random Generator] A function $\mathcal{G} : \{0, 1\}^\lambda \rightarrow \{0, 1\}^{\lambda+l(\lambda)}$ is a PRG if and only if

1. \mathcal{G} is computable in polynomial time;
2. $|\mathcal{G}(s)| = \lambda + l(\lambda)$ for all $s \in \{0, 1\}^\lambda$;
3. $\mathcal{G}(\mathcal{U}_\lambda) \approx_c \mathcal{U}_{\lambda+l(\lambda)}$.

◇

Theorem 7. *If \exists PRG with 1 bit of stretch, then \exists PRG with $l(\lambda)$ bits of stretch, with $l(\lambda) = \text{poly}(\lambda)$.* ◇

Proof of theorem 7. We'll prove this just for some fixed constant $l(\lambda) = l \in \mathbb{N}$.

First, let's look at the construction (fig. 3). We replicate our PRG \mathcal{G} with 1 bit stretch l times. The PRG \mathcal{G}^l that we define takes in input $s \in \{0, 1\}^\lambda$, computes $(s_1, b_1) = \mathcal{G}(s)$, where $s_1 \in \{0, 1\}^l$ and $b_1 \in \{0, 1\}$, outputs b_1 and feeds s_1 to the second copy of PRG \mathcal{G} , and so on until the l -th PRG.

To show that our construction is a PRG, we define l hybrids, with $\mathcal{H}_0^\lambda \equiv \mathcal{G}^l(\mathcal{U}_\lambda)$, where $\mathcal{G}^l : \{0, 1\}^\lambda \rightarrow \{0, 1\}^{\lambda+l}$ is our proposed construction, and \mathcal{H}_i^λ




Figure 3: Extending a PRG with 1 bit stretch to a PRG with l bit stretch.

takes $b_1, \dots, b_i \leftarrow \mathcal{S}\{0, 1\}$, $s_i \leftarrow \mathcal{S}\{0, 1\}^\lambda$, and outputs (b_1, \dots, b_i, s_l) , where $s_l \in \{0, 1\}^{\lambda+l-i}$ is $s_l = \mathcal{G}^{l-i}(s_i)$, *i.e.*, the output of our construction restricted to $l-i$ units.

\mathcal{H}_l^λ takes $b_1, \dots, b_l \leftarrow \mathcal{S}\{0, 1\}$ and $s_l \leftarrow \mathcal{S}\{0, 1\}^l$ and outputs (b_1, \dots, b_l, s_l) directly.

We need to show that $\mathcal{H}_i^\lambda \approx_c \mathcal{H}_{i+1}^\lambda$. To do so, fix some i . The only difference between the two hybrids is that s_{i+1}, b_{i+1} are pseudo random in \mathcal{H}_i^λ , and are truly random in $\mathcal{H}_{i+1}^\lambda$. All bits before them are truly random, all bits after are pseudo random.

Assume these two hybrids are distinguishable, then we can break the PRG. Consider the PPT function f_i defined by $f(s_{i+1}, b_{i+1}) = (b_1, \dots, b_l, s_l)$ such that $b_1, \dots, b_i \leftarrow \mathcal{S}\{0, 1\}$ and, for all $j \in [i+1, l]$ $(b_j, s_j) = \mathcal{G}(s_{j-1})$.

By the security of PRGs we have that $\mathcal{G}(\mathcal{U}_\lambda) \approx_c \mathcal{U}_{\lambda+1}$. By reduction, we also have that $f(\mathcal{G}(\mathcal{U}_\lambda)) \approx_c f(\mathcal{U}_{\lambda+1})$. Thus, $\mathcal{H}_i^\lambda \approx_c \mathcal{H}_{i+1}^\lambda$. \square

Acronyms

DL Discrete Log

CDH Computational Diffie-Hellman

DDH Decisional Diffie-Hellman

MAC Message Authentication Code

OTP One Time Pad

OWF One Way Function

PKC Public Key Cryptography

PPT Probabilistic Polynomial Time

PRG Pseudo Random Generator

RV random variable